

## FROM DESCRIPTIONS TO SOURCEABLE PARTS: CATALOG-GROUNDED RETRIEVAL FOR MECHANICAL DESIGN

Chancel Lavoie<sup>1</sup>, Levent Burak Kara<sup>1</sup>

<sup>1</sup>Carnegie Mellon University, Pittsburgh, PA

### ABSTRACT

*AI engineering agents are beginning to translate textual requirements into mechanical part and component descriptions. To support procurement, bill-of-material generation, and CAD insertion, these descriptions must be grounded in existing catalog parts with exact part numbers. Product-data APIs can attach live links, prices, datasheets, and CAD files after identifiers are known; catalog-grounded part-number retrieval is the upstream step that makes these APIs usable by design agents. We present an open-source Model Context Protocol server for headless navigation of McMaster-Carr that converts sufficiently detailed component descriptions into matching catalog part numbers. The system renders live catalog pages, extracts visible product-table schemas, uses a language model to map requested constraints onto live field names and values, and then performs deterministic row filtering. The output is the set of catalog parts that satisfy the grounded description. In final benchmark runs, the system recovered 250/250 exact part numbers across 25 catalog categories, returned the expected multi-part sets in 50/50 underspecified cases, and returned no part in 25/25 nonexistent cases. By making part sourcing callable by agents, the method enables downstream workflows in which assemblies are crafted from off-the-shelf parts and automatically linked to product data, prices, datasheets, CAD, and bills of materials.*

**Keywords:** AI engineering agents, design automation, part sourcing, model context protocol, catalog grounding, bill of materials, CAD retrieval

### 1. INTRODUCTION

AI engineering agents are increasingly used to generate mechanical part requirements, component selections, fixtures, subsystems, and assembly concepts from text. These outputs become actionable when they can be connected to existing parts that can be purchased, inspected, trusted, and inserted into CAD. Procurement, bill-of-materials generation, and CAD retrieval all depend on exact catalog identifiers, while generated requirements are typically expressed as natural-language component descriptions. This paper addresses the catalog-grounding step

between language-level engineering intent and sourceable mechanical parts.

McMaster-Carr is a widely used mechanical component catalog, and its Product Information API retrieves product information, live links, prices, images, CAD, and datasheets once a subscribed part number is known [1]. This identifier-centric interface is well suited to downstream product-data retrieval. Design agents, however, require an upstream capability: converting a textual component requirement into the catalog part numbers that satisfy it. A part-sourcing tool with this capability can connect generated component descriptions to live product data, CAD files, and bill-of-material entries.

Catalog-grounded part retrieval is challenging because industrial catalogs exhibit heterogeneous product schemas. A screw may be distinguished by material, thread size, length, drive style, package quantity, and finish. A switch may be distinguished by circuit, terminal count, maintained or momentary action, panel thickness, voltage, and terminal style. Conveyor rollers, springs, hinges, tapes, bearings, and filters each expose different discriminating fields. General retrieval therefore benefits from observing the live schema of the current catalog page and returning the complete set of rows that satisfy the description.

We present a dynamic catalog-grounding method exposed through a Model Context Protocol (MCP) server. Given a textual mechanical component description, the server returns the set of part numbers whose live catalog rows satisfy the grounded constraints. The system derives catalog search queries internally, renders live catalog pages, extracts product rows and visible column names, summarizes the live schema, aligns requested constraints with live fields and values, and filters rows programmatically. The language model serves as a schema-alignment component that produces grounded field-value matchers for the currently rendered catalog page. Once part numbers are recovered, existing product-data APIs can automatically attach live links, price information, datasheets, and CAD files to generated component or assembly descriptions.

The contribution of this work is an open-source MCP server that takes a textual mechanical component description and returns

the matching McMaster-Carr part numbers. The server packages headless rendered-page navigation, runtime product-schema extraction, language-model schema alignment, and deterministic row filtering behind a tool interface that can be called by engineering agents.

## 2. RELATED WORK

Our review focuses on three bodies of related work: tool protocols for connecting language models to external systems, McMaster-Carr product-data interfaces, and prior McMaster-Carr agent tooling.

### 2.1. Model Context Protocol for Engineering Tools

The Model Context Protocol standardizes how language-model applications connect to external context and tools [2]. MCP servers can expose resources, prompts, and tools; tools are the primitive used when a model needs to execute an action such as querying a database, calling an API, or driving an external system [3]. Our implementation uses MCP as the integration layer so that the same part-sourcing capability can be called by any compatible agent host.

MCP is useful here because part sourcing is not simply text generation. It requires a persistent browser session, rendered-page inspection, navigation, timeouts, and a structured result contract. These actions are better exposed as tools than hidden inside a prompt.

### 2.2. McMaster-Carr Product Data

McMaster-Carr’s Product Information API is available to approved customers and includes requests for product information, price, image retrieval, CAD retrieval, and datasheet retrieval once a subscribed part number is known [1]. This makes part-number discovery the critical upstream step. Our system is designed to fill that step and then hand part numbers to an approved product-data pipeline.

### 2.3. Existing McMaster-Carr Agent Tooling

A prior public repository, `mjbrown/mcmaster-agent`, exposes an MCP server for searching McMaster-Carr and looking up product information from part numbers [4]. Its documented MCP tools support product lookup from a known part number, broad search from a query string, and URL generation. The search tool returns product hits with part numbers and names, while the product lookup tool retrieves specification fields after a part number has already been selected. The repository also uses SeleniumBase UC mode, cookie persistence, and a visible browser window because its documentation reports that McMaster-Carr blocks headless browsers.

The distinction is therefore not simply implementation detail; it is the retrieval target. Prior tooling helps an agent search the catalog or enrich an already known part number. This work targets the upstream sourcing problem: given a textual component description, recover the catalog part numbers whose live rows satisfy the description. To do this, the server extracts rendered table columns, row attributes, filters, category links, and part-number rows from the current page; aligns the requested constraints to that live schema; and filters rows programmatically. The returned

object is a match set of part numbers, not a ranked search list or a product-data record for a known identifier. This difference is what makes the tool suitable for connecting generated mechanical requirements to downstream product-data, CAD, datasheet, and BOM retrieval.

## 3. SYSTEM DESIGN

### 3.1. MCP Server Surface

The package exposes a Python MCP server over stdio. The main user-facing tool is `mcmaster_find_exact_part`; the remaining tools expose the browser and extraction state needed for inspection, recovery, and agent-driven navigation when a first search page is insufficient. The public tools are:

- `mcmaster_find_exact_part`: resolve a detailed text description to the matching set of catalog part numbers;
- `mcmaster_extract_schema`: open or search a page and return dynamically extracted filters, table columns, row attributes, and part-number rows;
- `mcmaster_search`, `mcmaster_open`, `mcmaster_follow_link`, `mcmaster_current_page`, and `mcmaster_back`: provide stateful headless navigation through rendered pages;
- `mcmaster_find_parts`: broad search returning part numbers found during rendered search/category navigation; and
- `mcmaster_url`: generate product or search URLs without launching a browser.

These support tools are not separate research claims. They make the resolver usable by external agents: an agent can inspect the evidence behind a result, move through category pages, extract the current schema, and retry with better context without reimplementing browser automation. The implementation is packaged as `mcmaster-navigator-mcp`, requires Python 3.11 or newer, and uses SeleniumBase UC mode with headless Chrome. Model credentials are supplied at runtime rather than embedded in the package.

### 3.2. Headless Browser and Rendered-Page Extraction

McMaster-Carr pages are rendered web pages with product tables, category links, filters, and product detail pages. The navigator opens search URLs, auto-drills through category links when needed, and snapshots the rendered HTML. The extractor returns part numbers, product rows, row attributes, evidence text, navigable links, and schema objects containing filters, table columns, and part-number rows.

The browser is isolated by default with a temporary Chrome profile. The server serializes browser work through a single worker so browser state remains coherent across navigation calls. Browser logs are kept off MCP protocol stdout.

### 3.3. Dynamic Exact-Part Resolution

Figure 1 summarizes the exact-part resolver. The resolver derives catalog search roots and literal constraints from the description, searches the generated roots headlessly, extracts rows from

**Text requirement** → derived search roots and constraints → headless rendered pages → live rows and schema fields → LLM field/value matchers → deterministic row filtering → matching part-number set.

**FIGURE 1: Dynamic exact-part resolution.** The language model maps constraints to live page fields; Python filtering determines the returned set of matching part numbers.

product hits and schema tables, summarizes the live schema, asks the language model to map requested constraints to live fields and accepted values, and applies the resulting matchers with Python row filtering. If filtering conflicts with the live schema, the language model repairs the mapping from the extracted field/value summary. If one row remains, the resolver runs a final strict verification check against that row. If required constraints are missing or contradicted, the row is removed from the result set.

This design lets the system generalize across product families because the catalog schema is discovered at runtime. There is no static list of all possible attributes for screws, springs, switches, hinges, bearings, hose fittings, raw materials, or other categories.

### 3.4. Assembly-Level Use Case

The core problem is part sourcing. Assembly workflows become possible by treating each off-the-shelf line item as an independent exact-part resolution problem. A component requirement is sent to `mcmaster_find_exact_part`, and the returned matching set is carried forward. Downstream systems can retrieve product data, live links, price, image, CAD, and datasheet information for returned part numbers, preserve alternatives for later selection, or revise line items whose descriptions do not match any catalog row.

This line-item independence allows parallel sourcing agents. Compatibility checking can be layered above the resolver, but the core tool intentionally solves the narrower and measurable problem of finding the catalog part numbers that match stated requirements.

## 4. EXPERIMENT DESIGN

### 4.1. Exact Part Recovery

The exact recovery benchmark begins with known part numbers sampled from rendered McMaster-Carr search/category pages. Seed queries cover broad catalog areas including door hinges, toggle switches, welding clamps, socket head cap screws, air filter cartridges, drawer slides, compression springs, cartridge heaters, grease fittings, digital calipers, PVC tubing, timing belt pulleys, eye bolts, aluminum bar, wire rope clips, bearings, guide rails, pneumatic cylinders, double-sided tape, beakers, motors, and conveyor rollers.

For each sampled part number, the benchmark builds a textual description from the rendered product page or product-table row. The part number is withheld from the resolver. A case is counted correct if the returned set contains exactly one part number and that part number is the target. The final exact-recovery run used 250 cases, 25 catalog categories, gpt-5.4-mini, page budget 8, auto-drill depth 2, maximum schema rows 700, and two internally generated search roots.

### 4.2. Match-Set Size Checks

The near multi-match benchmark starts from successful exact cases, collects live schema rows around the target, then removes one discriminator from the description so that two to four parts are expected to remain. A case passes when the resolver returns the expected matching set, includes the original target part, and does not collapse the result to a single part.

The nonexistent benchmark starts from exact descriptions and appends impossible required constraints such as an unobtainable material, impossible color, and impossible size. The broad multi-match benchmark gives underspecified category queries such as “double sided tape,” “ball bearing,” “door hinges,” or “toggle switch.” These cases test whether the system returns the matching set instead of forcing a single best guess.

## 5. EXPERIMENTAL RESULTS

Table 1 reports the final benchmark runs. The exact-recovery benchmark returned exactly one part number for all 250 cases, and that part number was the target in every case. Median exact-case latency was 54.448 s, p90 latency was 85.960 s, p95 latency was 161.763 s, and maximum latency was 466.125 s. Token use averaged 7,449.96 tokens per case, with median 6,962.5, p90 10,563.9, p95 11,183.6, and maximum 18,300.

The near multi-match benchmark returned candidate counts from 2 to 4, with mean 2.24. The expected matching set coverage was 1.0 in all cases. The nonexistent benchmark returned empty part-number sets for all 25 cases. The broad multi-match benchmark returned multiple part numbers for all 25 cases. Returned candidate counts ranged from 2 to 456, with median 86 and mean 128.88.

The absence of misses across these categories is evidence that runtime schema extraction can handle substantially different product families. It is not proof of complete coverage of every McMaster-Carr page or every possible human description. The strongest defensible claim is that dynamic schema extraction and deterministic filtering can recover exact part numbers across diverse catalog families when descriptions contain sufficient identifying information.

## 6. CONCLUSION

This work turns a practical part-sourcing bottleneck into a measurable tool interface. A design agent can describe a required component, but a downstream product information API needs a part number. The proposed MCP server bridges that gap by headlessly navigating McMaster-Carr, extracting live product schemas, grounding textual constraints to live fields, and filtering rows deterministically. In final benchmark runs, the system recovered 250/250 exact part numbers across 25 catalog categories, returned expected multi-part sets when requirements were under-specified, and returned no part when requirements could not be satisfied. Reliable part sourcing can let agents compose off-the-shelf components into larger designs while automatically attaching BOM entries, CAD, datasheets, prices, live links, and product data.

**TABLE 1: Final benchmark results. Tokens are total recorded model tokens for each final run.**

| Evaluation                 | Cases | Expected count | Pass criterion                 | Passed  | Pass rate | Tokens    | Mean s/case |
|----------------------------|-------|----------------|--------------------------------|---------|-----------|-----------|-------------|
| Exact single-part recovery | 250   | 1              | Target returned as only part   | 250/250 | 100%      | 1,862,490 | 71.038      |
| Near multi-match           | 25    | 2–4            | Expected matching set returned | 25/25   | 100%      | 187,701   | 64.325      |
| Nonexistent requirements   | 25    | 0              | No part returned               | 25/25   | 100%      | 268,293   | 75.348      |
| Broad multi-match          | 25    | > 1            | Multiple parts returned        | 25/25   | 100%      | 90,950    | 50.654      |
| Total final reported set   | 325   | mixed          | Task-specific                  | 325/325 | 100%      | 2,409,434 | –           |

**TABLE 2: Exact-recovery category coverage. Each category had 100% single-part recovery in the final run.**

| Category                  | Cases | Single |
|---------------------------|-------|--------|
| Adhesives                 | 8     | 8      |
| Bearings                  | 8     | 8      |
| Building & Grounds        | 20    | 20     |
| Conveying                 | 8     | 8      |
| Electrical & Lighting     | 6     | 6      |
| Fabricating               | 19    | 19     |
| Fastening & Joining       | 20    | 20     |
| Filtering                 | 21    | 21     |
| Furniture & Storage       | 14    | 14     |
| Hand Tools                | 8     | 8      |
| Hardware                  | 8     | 8      |
| Heating & Cooling         | 8     | 8      |
| Lab Supplies              | 8     | 8      |
| Linear Motion             | 7     | 7      |
| Lubricating               | 8     | 8      |
| Measuring & Inspecting    | 8     | 8      |
| Motors                    | 8     | 8      |
| Pipe/Tubing/Hose/Fittings | 8     | 8      |
| Plumbing & Janitorial     | 8     | 8      |
| Pneumatics                | 8     | 8      |
| Power Transmission        | 8     | 8      |
| Pulling & Lifting         | 8     | 8      |
| Raw Materials             | 8     | 8      |
| Shipping                  | 8     | 8      |
| Suspending                | 7     | 7      |

## 7. DISCUSSION

### 7.1. Why Dynamic Schemas Matter

The benchmark covers product families whose discriminating attributes differ substantially. A static parser with a fixed list of fields would need to know domain-specific names and value formats for hinges, switches, bearings, adhesives, filters, raw materials, pneumatic cylinders, motors, and conveyor rollers. The dynamic approach avoids that requirement. The system only needs a general row representation and a way to align user constraints with fields currently visible on the page.

### 7.2. Why Returning the Set Matters

Procurement workflows need faithful part sets. Returning a single best guess for “ball bearing” is worse than returning the matching catalog rows, because the description does not specify shaft diameter, housing diameter, width, seal type, load capacity, or material. The multi-match and zero-match experiments show

that the resolver can return all matching parts, or none, without false precision.

### 7.3. Integration with Product-Data APIs

Once part numbers are known, product-data retrieval is a conventional API task. The McMaster-Carr Product Information API includes endpoints for product information, price, images, CAD, and datasheets [1]. This separation lets the proposed MCP server remain focused on the hard upstream retrieval problem: finding the part numbers from text and navigation.

## 8. LIMITATIONS AND FUTURE WORK

The benchmark descriptions were generated from rendered catalog rows and product pages. They are specific and often include labels such as family, group, selected option, and attribute names. This is a valid test of schema grounding and exact retrieval, but it is not the same as testing noisy human design notes.

The system depends on rendered website structure. If McMaster-Carr changes page layout, table markup, anti-automation behavior, or part-number presentation, extraction may require maintenance.

The system is not affiliated with McMaster-Carr and should be used in accordance with applicable terms and approved API access. The official product information API requires approved customer access and subscription to product information.

Latency is dominated by headless browsing. The mean exact-recovery latency in the final run was 71.038 seconds per case. Parallel line-item sourcing can reduce assembly-level wall-clock time, but single-line resolution remains slow relative to a pure API call.

Future work should evaluate noisier human-written requirements, demonstrate full assembly-level sourcing with compatibility checks across line items, and measure robustness over time as catalog pages change.

## REFERENCES

- [1] McMaster-Carr. “McMaster-Carr Product Information API.” McMaster-Carr (2026). Accessed 2026-05-11, URL <https://www.mcmaster.com/help/api/>.
- [2] Model Context Protocol. “Model Context Protocol Specification.” Model Context Protocol (2024). Accessed 2026-05-11, URL <https://modelcontextprotocol.io/specification/2024-11-05/index>.
- [3] Model Context Protocol. “Tools.” Model Context Protocol (2026). Accessed 2026-05-11, URL <https://modelcontextprotocol.io/specification/draft/server/tools>.
- [4] Braun, Michael J. “McMaster-Carr MCP Server.” GitHub (2026). Accessed 2026-05-11, URL <https://github.com/mjbraun/mcmaster-agent>.